

# Ben Roth Week 6 - Jupyter Notebook

March 22, 2020

## 1 Assignment is at the bottom!

```
In [1]: from sklearn.linear_model import LogisticRegression
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np


from pylab import rcParams
rcParams['figure.figsize'] = 20, 10

from sklearn.linear_model import LogisticRegression as Model

In [2]: y = np.concatenate([np.zeros(10), np.ones(10)])
x = np.linspace(0, 10, len(y))

In [3]: plt.scatter(x, y, c=y)

Out[3]: <matplotlib.collections.PathCollection at 0x1a1fff80f0>
```



```
In [4]: model = LogisticRegression()
```


```
In [5]: model.fit(x.reshape(-1, 1),y)
```

```
/Users/Broth/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning
```

```
Out[5]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='warn',
    n_jobs=None, penalty='l2', random_state=None, solver='warn',
    tol=0.0001, verbose=0, warm_start=False)
```

```
In [6]: plt.scatter(x,y, c=y)
    plt.plot(x, model.predict_proba(x.reshape(-1, 1))[:,1])
```

```
Out[6]: [<matplotlib.lines.Line2D at 0x1a200e3358>]
```




```
In [7]: b, b0 = model.coef_, model.intercept_
        model.coef_, model.intercept_
```

```
Out[7]: (array([[0.47990097]]), array([-1.84794266]))
```

```
In [8]: plt.plot(x, 1/(1+np.exp(-x)))
```

```
Out[8]: [<matplotlib.lines.Line2D at 0x1a20588128>]
```




```
In [9]: b
```

```
Out[9]: array([[0.47990097]])
```

```
In [10]: plt.plot(x, 1/(1+np.exp(-(b[0]*x +b0))))
```

```
Out[10]: [matplotlib.lines.Line2D at 0x1a2073a080>]
```



```
In [11]: from mpl_toolkits.mplot3d import Axes3D # noqa: F401 unused import
```

```
import matplotlib.pyplot as plt
from matplotlib import cm
from matplotlib.ticker import LinearLocator, FormatStrFormatter
import numpy as np
```


```
fig = plt.figure()
ax = fig.gca(projection='3d')
```

```
# Make data.
X = np.arange(-10, 10, 0.25)
Y = np.arange(-10, 10, 0.25)
X, Y = np.meshgrid(X, Y)
```

```

R = np.sqrt(X**2 + Y**2)
Z = 1/(1+np.exp(-(b[0]*X +b[0]*Y +b0)))
surf = ax.plot_surface(X, Y, Z, cmap=cm.coolwarm,
                      linewidth=0, antialiased=False)

```



In [12]: X

```

Out[12]: array([[ -10.   ,  -9.75,  -9.5  , ...,   9.25,   9.5  ,   9.75],
                [ -10.   ,  -9.75,  -9.5  , ...,   9.25,   9.5  ,   9.75],
                [ -10.   ,  -9.75,  -9.5  , ...,   9.25,   9.5  ,   9.75],
                ...,
                [ -10.   ,  -9.75,  -9.5  , ...,   9.25,   9.5  ,   9.75],
                [ -10.   ,  -9.75,  -9.5  , ...,   9.25,   9.5  ,   9.75],
                [ -10.   ,  -9.75,  -9.5  , ...,   9.25,   9.5  ,   9.75]])

```

In [13]: Y

```

Out[13]: array([[ -10.   , -10.   , -10.   , ..., -10.   , -10.   , -10.   ],
                [ -9.75,  -9.75,  -9.75, ..., -9.75,  -9.75,  -9.75],
                [ -9.5  ,  -9.5  ,  -9.5  , ..., -9.5  ,  -9.5  ,  -9.5  ],
                ...,
                [  9.25,   9.25,   9.25, ...,   9.25,   9.25,   9.25],
                [  9.5  ,   9.5  ,   9.5  , ...,   9.5  ,   9.5  ,   9.5  ],
                [  9.75,   9.75,   9.75, ...,   9.75,   9.75,   9.75]])

```

What if the data doesn't really fit this pattern?


```

In [14]: y = np.concatenate([np.zeros(10), np.ones(10), np.zeros(10)])
         x = np.linspace(0, 10, len(y))

```

```
In [15]: plt.scatter(x,y, c=y)
```

```
Out[15]: <matplotlib.collections.PathCollection at 0x1a20ae6f28>
```




```
In [16]: model.fit(x.reshape(-1, 1),y)
```

```
/Users/Brth/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning
```

```
Out[16]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='warn',
    n_jobs=None, penalty='l2', random_state=None, solver='warn',
    tol=0.0001, verbose=0, warm_start=False)
```

```
In [17]: plt.scatter(x,y)
plt.plot(x, model.predict_proba(x.reshape(-1, 1)))
```

```
Out[17]: [<matplotlib.lines.Line2D at 0x1a20ae7668>,
    <matplotlib.lines.Line2D at 0x1a20a93860>]
```



```
In [18]: model1 = LogisticRegression()
         model1.fit(x[:15].reshape(-1, 1), y[:15])
```

```
/Users/Broth/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning
```

```
Out[18]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, max_iter=100, multi_class='warn',
                             n_jobs=None, penalty='l2', random_state=None, solver='warn',
                             tol=0.0001, verbose=0, warm_start=False)
```


```
In [19]: model2 = LogisticRegression()
         model2.fit(x[15:].reshape(-1, 1), y[15:])
```

```
/Users/Broth/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning
```

```
Out[19]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, max_iter=100, multi_class='warn',
                             n_jobs=None, penalty='l2', random_state=None, solver='warn',
                             tol=0.0001, verbose=0, warm_start=False)
```

```
In [20]: plt.scatter(x, y, c=y)
         plt.plot(x, model1.predict_proba(x.reshape(-1, 1))[:, 1] * model2.predict_proba(x.reshape(-1, 1))[:, 1], c='b')
```

Out[20]: [<matplotlib.lines.Line2D at 0x1a20894e48>]



```
In [21]: df = pd.read_csv('../data/adult.data', index_col=False)
golden = pd.read_csv('../data/adult.test', index_col=False)
```

```
In [ ]:
```

```
In [22]: from sklearn import preprocessing
```

```
enc = preprocessing.OrdinalEncoder()
```

```
In [23]: transform_columns = ['sex', 'workclass', 'education', 'marital-status',
                              'occupation', 'relationship', 'race', 'sex',
                              'native-country', 'salary']
```

```
In [24]: x = df.copy()
```

```
x[transform_columns] = enc.fit_transform(df[transform_columns])
```

```
golden['salary'] = golden.salary.replace(' <=50K.', ' <=50K').replace(' >50K.', ' >50K')
xt = golden.copy()
```

```
xt[transform_columns] = enc.transform(golden[transform_columns])
```

```
In [25]: df.salary.unique()
```

```
Out[25]: array([' <=50K', ' >50K'], dtype=object)
```



```
In [26]: golden.salary.replace(' <=50K.', ' <=50K').replace(' >50K.', ' >50K').unique()
```

```
Out[26]: array([' <=50K', ' >50K'], dtype=object)
```

```
In [27]: model.fit(preprocessing.scale(x.drop('salary', axis=1)), x.salary)
```

```
/Users/Broth/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:1: DataConversionWarn
    """Entry point for launching an IPython kernel.
/Users/Broth/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:433: Future
    FutureWarning)
```

```
Out[27]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='warn',
    n_jobs=None, penalty='l2', random_state=None, solver='warn',
    tol=0.0001, verbose=0, warm_start=False)
```

```
In [28]: pred = model.predict(preprocessing.scale(x.drop('salary', axis=1)))
    pred_test = model.predict(preprocessing.scale(xt.drop('salary', axis=1)))
```

```
/Users/Broth/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:1: DataConversionWarn
    """Entry point for launching an IPython kernel.
/Users/Broth/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:2: DataConversionWarn
```

```
In [29]: x.head()
```

```
Out[29]:
```

	age	workclass	fnlwgt	education	education-num	marital-status	\
0	39	7.0	77516	9.0	13	4.0	
1	50	6.0	83311	9.0	13	2.0	
2	38	4.0	215646	11.0	9	0.0	
3	53	4.0	234721	1.0	7	2.0	
4	28	4.0	338409	9.0	13	2.0	

	occupation	relationship	race	sex	capital-gain	capital-loss	\
0	1.0	1.0	4.0	1.0	2174	0	
1	4.0	0.0	4.0	1.0	0	0	
2	6.0	1.0	4.0	1.0	0	0	
3	6.0	0.0	2.0	1.0	0	0	
4	10.0	5.0	2.0	0.0	0	0	

	hours-per-week	native-country	salary
0	40	39.0	0.0
1	13	39.0	0.0
2	40	39.0	0.0
3	40	39.0	0.0
4	40	5.0	0.0

```
In [30]: from sklearn.metrics import (
        accuracy_score,
        classification_report,
        confusion_matrix, auc, roc_curve
    )

In [31]: accuracy_score(x.salary, pred)

Out[31]: 0.8250667977027732

In [32]: confusion_matrix(x.salary, pred)

Out[32]: array([[23300, 1420],
               [ 4276, 3565]])

In [33]: print(classification_report(x.salary, pred))
```

	precision	recall	f1-score	support
0.0	0.84	0.94	0.89	24720
1.0	0.72	0.45	0.56	7841
micro avg	0.83	0.83	0.83	32561
macro avg	0.78	0.70	0.72	32561
weighted avg	0.81	0.83	0.81	32561

```
In [34]: print(classification_report(xt.salary, pred_test))
```

	precision	recall	f1-score	support
0.0	0.85	0.94	0.89	12435
1.0	0.70	0.45	0.55	3846
micro avg	0.82	0.82	0.82	16281
macro avg	0.77	0.69	0.72	16281
weighted avg	0.81	0.82	0.81	16281

## 2 Assignment

**2.1 1. Use your own dataset (create a train and a test set) and build 2 models: Logistic Regression and Decision Tree (shallow). Compare the test results.**

```
In [35]: import random
        random.seed(1234)
```

```

df = pd.read_stata('../data/Titanic.dta')

df.head(5)

```

Out [35]:

	pclass	survived	name	sex
0	1	1	Allen, Miss. Elisabeth Walton	female
1	1	1	Allison, Master. Hudson Trevor	male
2	1	0	Allison, Miss. Helen Loraine	female
3	1	0	Allison, Mr. Hudson Joshua Creighton	male
4	1	0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	female

	age	sibsp	parch	ticket	fare	cabin	embarked
0	29.00	0	0	24160	211.337494	B5	S
1	0.92	1	2	113781	151.550003	C22 C26	S
2	2.00	1	2	113781	151.550003	C22 C26	S
3	30.00	1	2	113781	151.550003	C22 C26	S
4	25.00	1	2	113781	151.550003	C22 C26	S

```

In [36]: df = df.drop(['name', 'ticket', 'cabin'], axis = 1)

from sklearn.preprocessing import LabelEncoder

str_col = df.select_dtypes('object').columns

for col in str_col:
    df[col] = LabelEncoder().fit_transform(df[col].values)

In [37]: df.isna().sum()

Out [37]: pclass      0
survived      0
sex           0
age          263
sibsp         0
parch         0
fare          1
embarked      0
dtype: int64

In [38]: df = df.dropna()

In [39]: X = df.drop('survived', axis = 1)
y = df['survived']

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)

In [40]: from sklearn.tree import DecisionTreeClassifier

```

```

logit = LogisticRegression()

logit.fit(X_train, y_train)

tree = DecisionTreeClassifier(max_depth = 3)

tree.fit(X_train, y_train)

/Users/Bruth/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning
FutureWarning)

```

```

Out[40]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=3,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                                splitter='best')

```

```

In [41]: log_pred = logit.predict(X_test)
         tree1_pred = tree.predict(X_test)

         print('The classification report for the logistic regression model is:\n', classification_report(X_test, log_pred))

         print('\nThe classification report for the shallow decision tree is:\n', classification_report(X_test, tree1_pred))

```

The classification report for the logistic regression model is:

	precision	recall	f1-score	support
0	0.88	0.82	0.85	222
1	0.71	0.80	0.76	123
micro avg	0.81	0.81	0.81	345
macro avg	0.80	0.81	0.80	345
weighted avg	0.82	0.81	0.82	345

The classification report for the shallow decision tree is:

	precision	recall	f1-score	support
0	0.89	0.81	0.85	226
1	0.69	0.81	0.74	119
micro avg	0.81	0.81	0.81	345
macro avg	0.79	0.81	0.80	345
weighted avg	0.82	0.81	0.81	345

Looking at the two classification reports, the two models appear comparable, although the logit model performs slightly better.

## 2.2 2. Repeat 1. but let the Decision Tree be much deeper to allow over-fitting. Compare the two models' test results again, does the Logistic Regression have an improvement due to a lower variance

```
In [42]: deep_tree = DecisionTreeClassifier(max_depth = 20)
         deep_tree.fit(X_train, y_train)
         deep_tree_preds = deep_tree.predict(X_test)
         print('The classification report for the deep decision tree is:\n', classification_report(deep_tree_preds, y_test))

         print('\nThe classification report for the logistic regression model is:\n', classification_report(logit_model_preds, y_test))
```

The classification report for the deep decision tree is:

	precision	recall	f1-score	support
0	0.81	0.82	0.81	206
1	0.72	0.71	0.72	139
micro avg	0.77	0.77	0.77	345
macro avg	0.77	0.76	0.76	345
weighted avg	0.77	0.77	0.77	345

The classification report for the logistic regression model is:

	precision	recall	f1-score	support
0	0.88	0.82	0.85	222
1	0.71	0.80	0.76	123
micro avg	0.81	0.81	0.81	345
macro avg	0.80	0.81	0.80	345
weighted avg	0.82	0.81	0.82	345

Comparing the deeper decision tree to the logistic regression model, the LR model clearly performs better. As suggested in the question, the deeper decision tree is over-fit